

Celery - Distributed Task Queue



Mike DeLaurentis

May 21, 2013

What is it?

- Distributed task queue
- Can run tasks in one process,
- ... or multiple processes on one machine
- ... or multiple machines



Setting up

- Installs easily with pip
- Needs a “broker” (like RabbitMQ, Redis, or relational DB) to hold tasks
- May need a “backend” if you want direct access to results
- Config is easy,
 - either with a simple file
 - or programmatically

```
from celery import Celery

celery = Celery(
    'pug_demo',
    broker='redis://localhost/1',
    backend='redis')
```

Writing tasks

- Define tasks as decorated functions
- Takes serializable args
- Can return some result
- ... or store it elsewhere

```
@celery.task
def say(*msg):
    print("".join(msg))

@celery.task
def increment(x):
    return x + 1
```

Running workers

- Start a worker by running 'celery'
- Identify module with tasks using --app
- Specify number of threads with --concurrency

```
delaurentis@Mikes-MacBook-Air ~/src/talks/2013-05-21-Celery[master*]
$ celery worker --app pugdemo --concurrency 4

----- celery@Mikes-MacBook-Air.local v3.0.18 (Chiastic Slide)
----- **** -----
-- * *** * -- Darwin-11.4.2-x86_64-i386-64bit
-- * - **** --
- ** ----- [config]
- ** ----- .> broker:      redis://localhost:6379/1
- ** ----- .> app:        pugdemo:0x10148e150
- ** ----- .> concurrency: 4 (processes)
- *** --- * --- .> events:    OFF (enable -E to monitor this worker)
-- ****** --
-- ***** --- [queues]
----- .> celery:      exchange:celery(direct) binding:celery

[2013-05-19 22:17:32,780: WARNING/MainProcess] celery@Mikes-MacBook-Air.local ready.
```



Submitting tasks

```
# Run synchronously
say('I am synchronous')

# Run asynchronously.
# result.get() waits for task to finish
result = increment.delay((5))
print("5 + 1 is", result.get())

# Make 'subtask' to run later with args
# Args at call-time are prepended
ask = say.s("?")
task = ask("How are you")
task.get()
```

- Calling directly will run in process
- Call `delay(*args)` to put on queue
- Call `s(*args)` to make task object



Creating workflows

- Build a workflow programatically, out of:
 - chain - List of tasks that must be run in order
 - group - Tasks that can be run simultaneously
 - chord - Group with callback
 - chunks - Split a task with a long list of args into smaller tasks
- Compose these elements to make complex workflows



Immutable workflow

```
prep = group(  
    say.si("slice bread"),  
    say.si("slice cheese"),  
    say.si("put butter in pan"))  
  
grilledcheese = chain(  
    prep |  
    say.si("turn on burner") |  
    say.si("assemble") |  
    say.si("cook one side") |  
    say.si("cook other side"))  
  
grilledcheese()
```



Mutable workflow

```
add_three_and_square = chain(
    increment.s() |
    increment.s() |
    increment.s() |
    square.s())

print("Add three and square: ", add_three_and_square)
res = add_three_and_square(5)
print("(5 + 3) ^ 2 =", res.get())
```



Thanks!

